

verl: an flexible and efficient RL framework for LLMs

Haibin Lin, Bytedance MLSys & LMSYS.org



Reinforcement learning

- Supervised fine-tuning
 - Learning from labeled examples
- Reinforcement learning
 - Optimization based on rewards
 - Preference alignment
 - Human feedbacks
 - Reasoning with automated feedbacks
 - Coding: unit-tests
 - Math: ground truth graders
 - Agentic tasks: operator, deep research

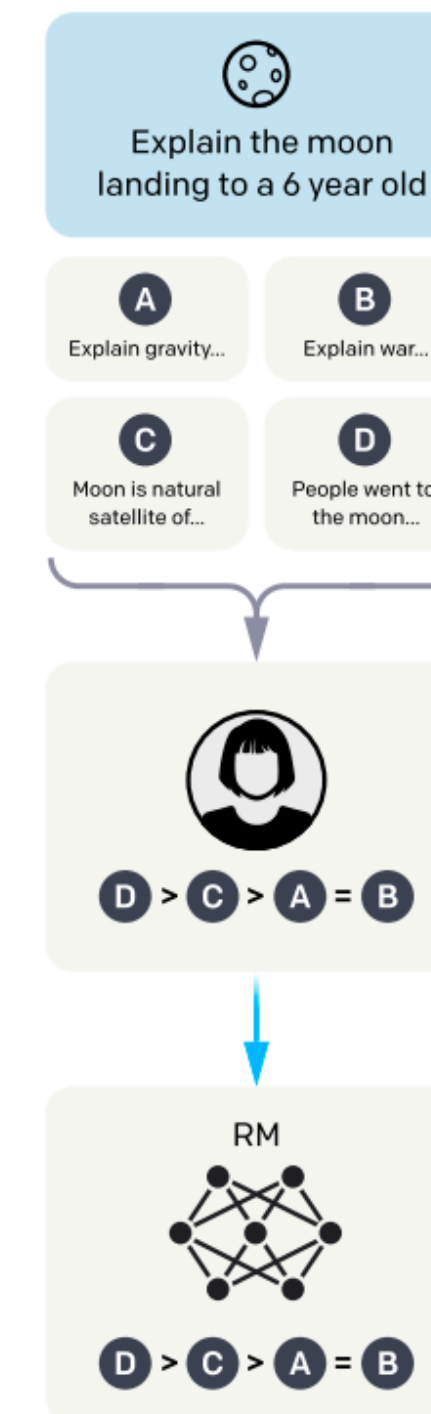
Step 2

Collect comparison data, and train a reward model.

A prompt and several model outputs are sampled.

A labeler ranks the outputs from best to worst.

This data is used to train our reward model.



Step 3

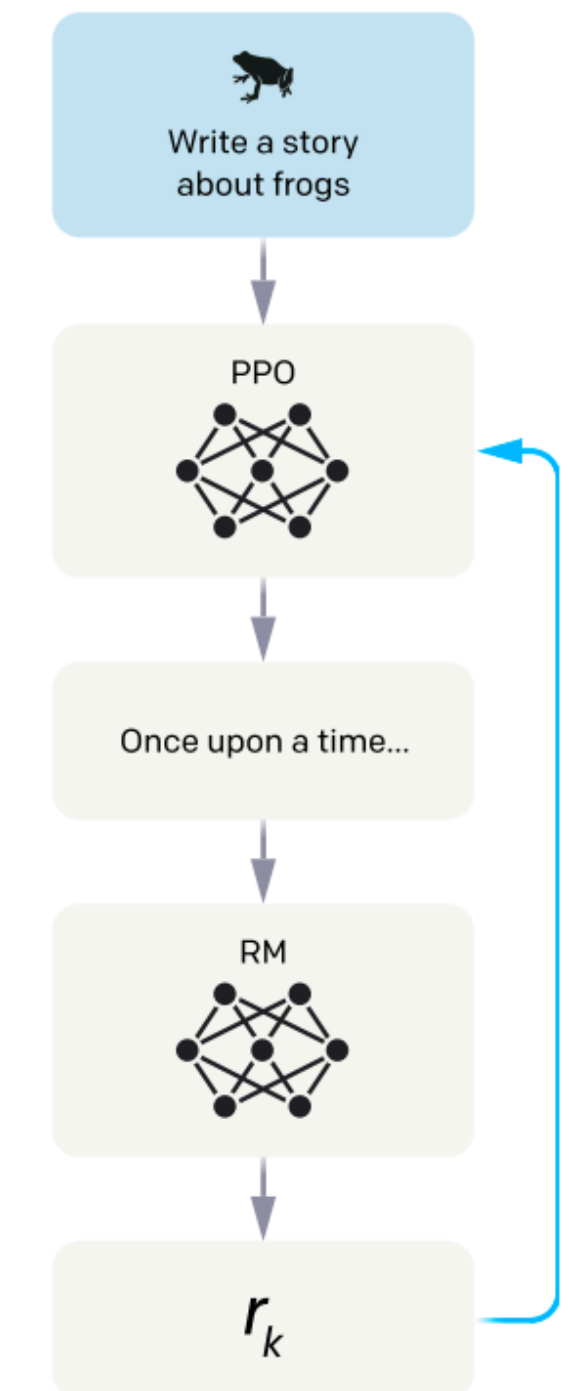
Optimize a policy against the reward model using reinforcement learning.

A new prompt is sampled from the dataset.

The policy generates an output.

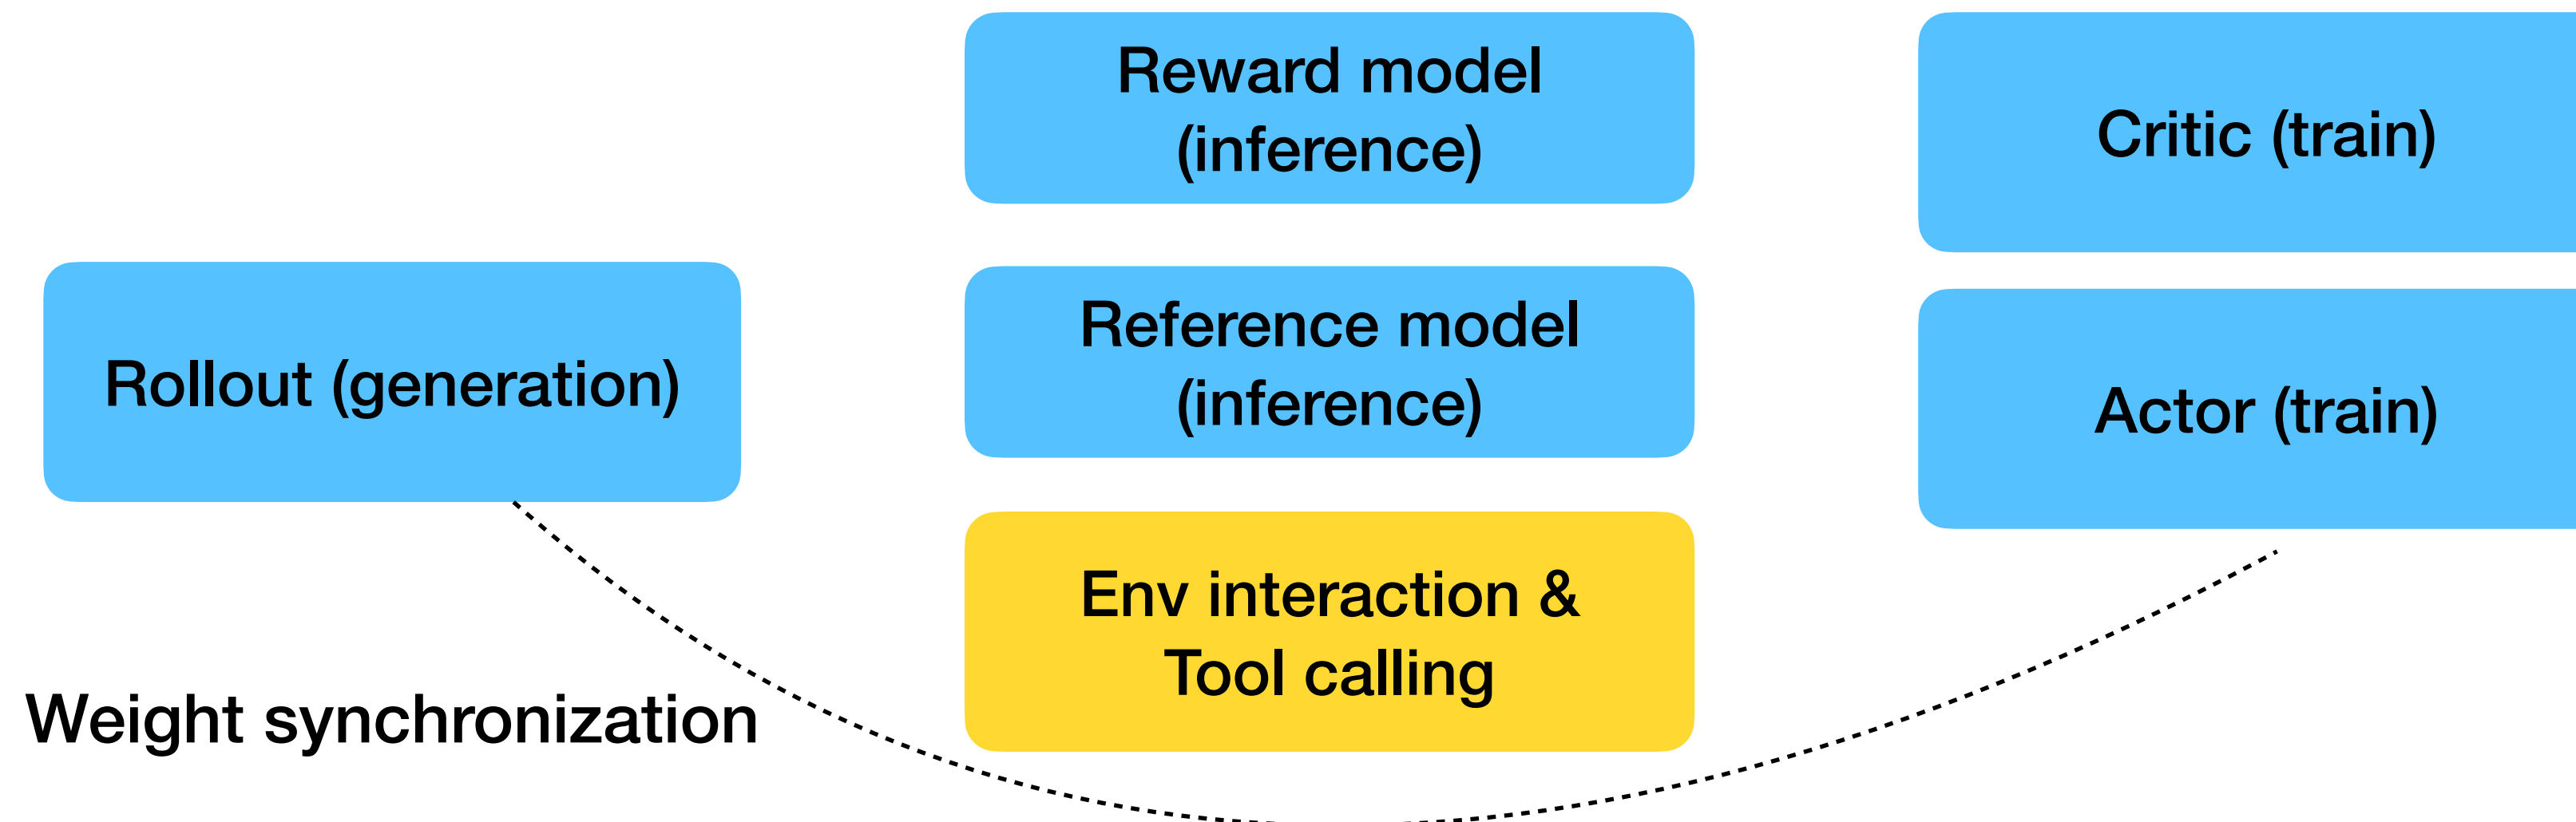

The reward model calculates a reward for the output.

The reward is used to update the policy using PPO.



Infra challenges for RL on LLMs

- the need for nD parallelisms (Megatron-LM)
 - Growing model size: llama 70b, Deepseek 671B
 - Growing sequence length: 8k -> 1M
- the need for programming abstractions

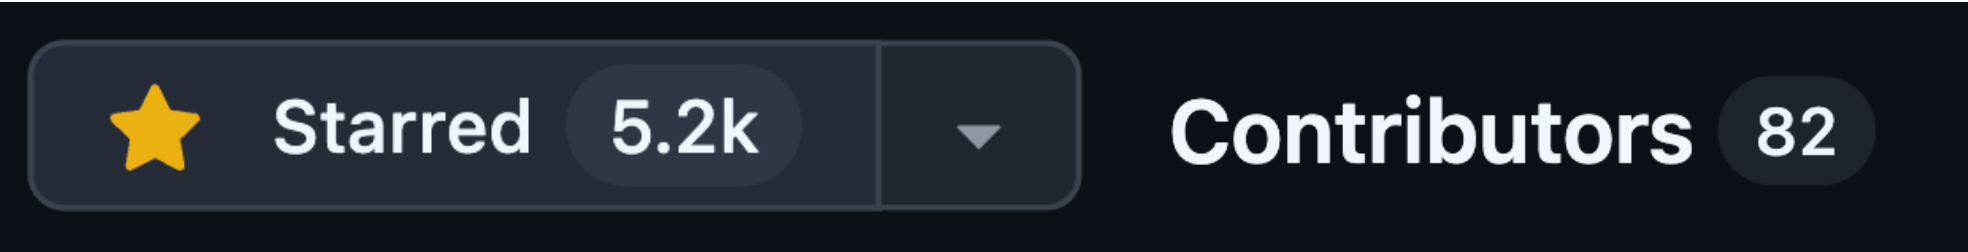


verl: open framework for RL on LLMs



History & community

- Developed & adopted internally since 2023/9 for research*, open sourced on 2024/10
- Reinforcement learning at Bytedance
 - Reasoning: O1-level performance on math benchmarks
 - RLHF, Image generation, music generation
 - Desktop operator, coding assistant...
- Users and contributors:




★ Starred 5.2k ▾ Contributors 82

 - PKU, THU, UIUC, UCB, UCLA, HKU, Stanford, Northwestern, MIT...
 - Amazon, NVIDIA, LMSys, Alibaba, StepFun, Anyscale, OpenHands, ...

*HybridFlow: A Flexible and Efficient RLHF Framework (Eurosyst 25')

verl: open framework for RL on LLMs



Features...

- RL recipes: PPO, GRPO, RLOO, reinforce++, DAPO
- Transformers integration: deepseek, llama, qwen, gemma, etc
- Inference engine: vllm, sglang**
- Distributed training engine: FSDP, Megatron
- System optimizations: seq packing, seq parallelism, fused entropy kernels
- Hardware support: NVIDIA GPU, AMD GPU*, Huawei NPU**
- Hybrid controller programming

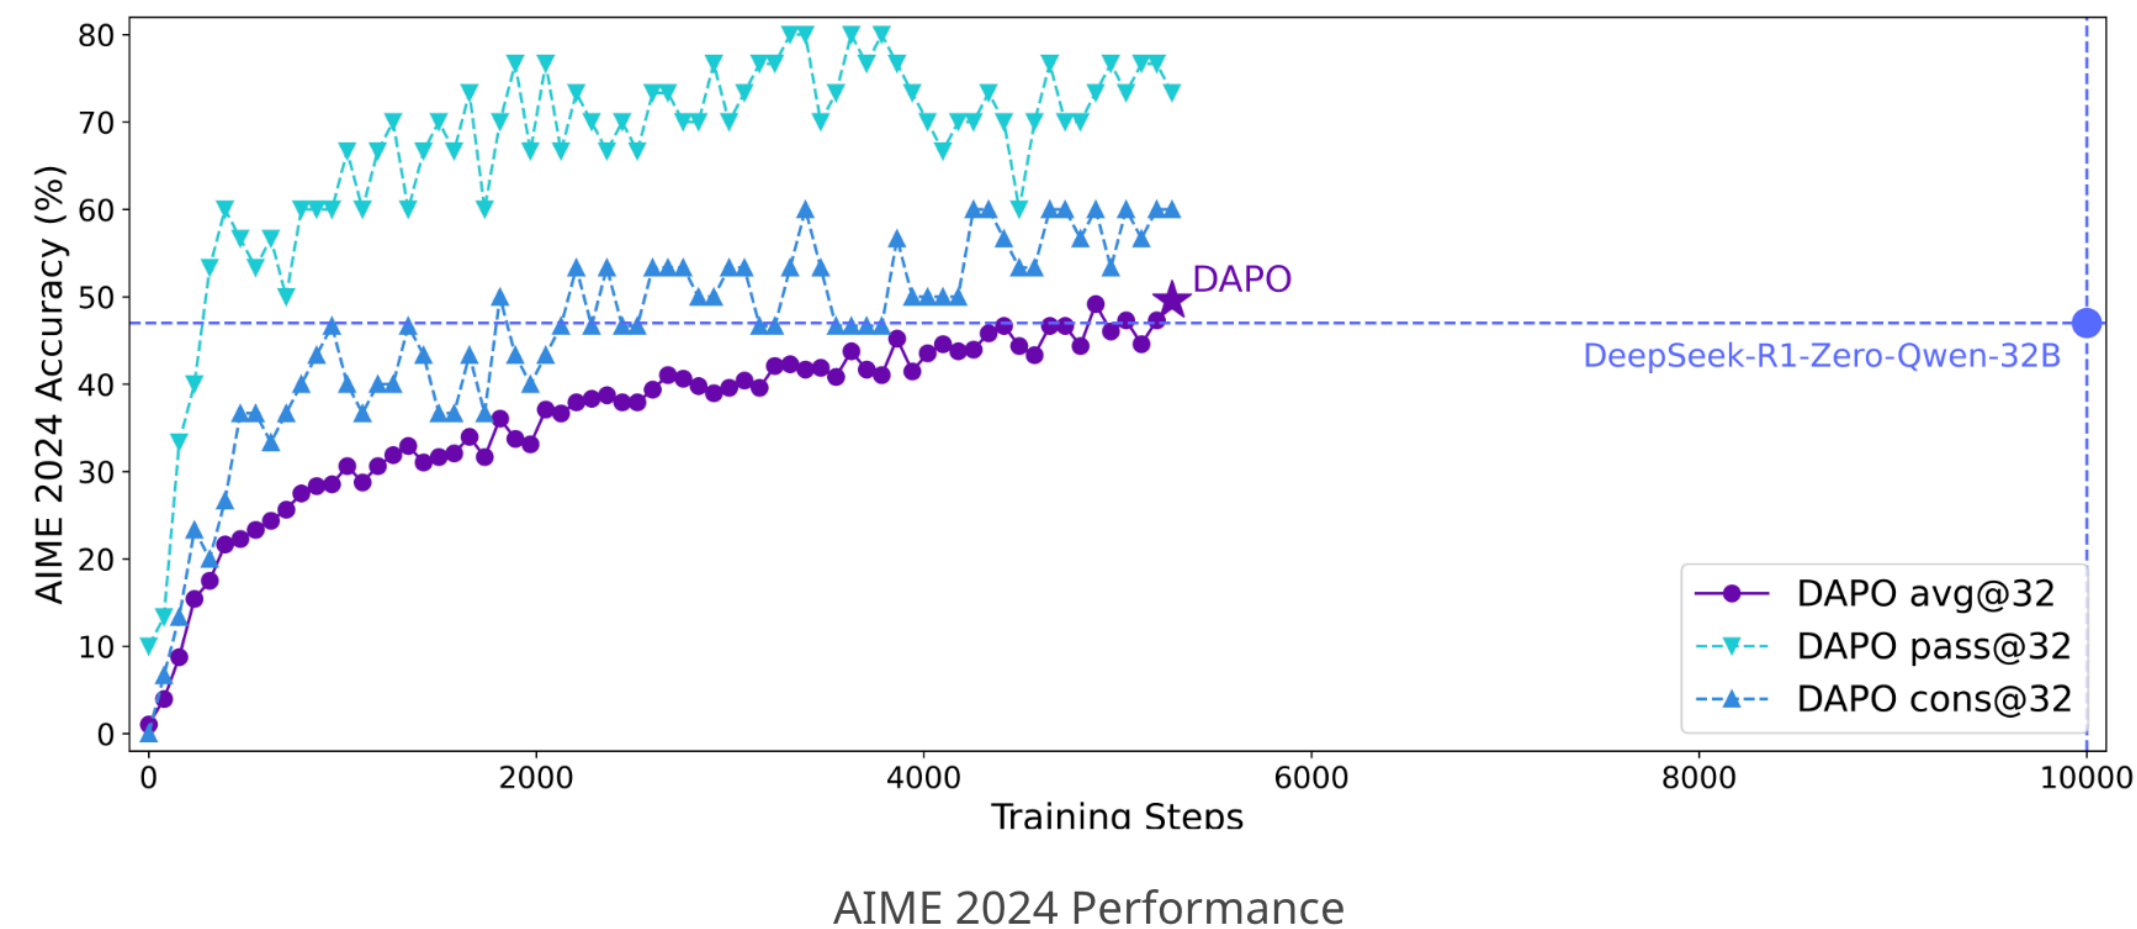
** experimental

verl for LLM RL: why is it special?

flexible, efficient, battle-tested



- **DAPO algorithm:** improvements on top of GRPO
- beats DeepSeek-R1-zero-32B with fewer steps
- Fully open source recipe (dataset, code, logs, model)



*DAPO: an Open-Source LLM Reinforcement Learning System at Scale

- **hybrid controller programming model***
- program distributed RL algorithm like a single threaded program

```
# Initialize cost model by reusing the RewardWorker
cost = RewardWorker(cost_config, resource_pool)
... # omit other models initialization
algo_type = "Safe-RLHF" # specify different RLHF numerical computation.
# Examples of PPO and Safe-RLHF
for (prompts, pretrain_batch) in dataloader:
    # Stage 1: Generate responses
    batch = actor.generate_sequences(prompts)
    batch = actor.generate_sequences(prompts, do_sample=False)
    # Stage 2: Prepare experience
    batch = critic.compute_values(batch)
    batch = reference.compute_log_prob(batch)
    batch = reward.compute_reward(batch)
    batch = cost.compute_cost(batch)
    batch = compute_advantages(batch, algo_type)
    # Stage 3: Actor and critic training
    critic_metrics = critic.update_critic(batch, loss_func=algo_type)
    pretrain_loss = actor.compute_loss(pretrain_batch)
    batch["pretrain_loss"] = pretrain_loss
    actor_metrics = actor.update_actor(batch, loss_func=algo_type)
```

Annotations:

- `batch = actor.generate_sequences(prompts, do_sample=False)` is added for ReMax
- `batch = critic.compute_values(batch)` is added for ReMax
- `batch = cost.compute_cost(batch)` is added for Safe-RLHF
- `critic_metrics = critic.update_critic(batch, loss_func=algo_type)` is added for Safe-RLHF
- `batch["pretrain_loss"] = pretrain_loss` is added for Safe-RLHF

*HybridFlow: A Flexible and Efficient RLHF Framework (Eurosys 25')

verl Roadmap



- megatron v0.11 for deepseek-v3 MOE
- multi-turn optimizations
- agentic environment & tool calling
- stable sglang integration
- stable hardware support: AMD & NPU
- contributions are welcome!!

Awesome work using verl

- [TinyZero](#): a reproduction of DeepSeek R1 Zero recipe for reasoning tasks
- [RAGEN](#): a general-purpose reasoning agent training framework
- [deepscaler](#): iterative context scaling with GRPO
- [Easy-R1](#): Multi-modal RL training framework
- [self-rewarding-reasoning-LLM](#): self-rewarding and correction with generative reward model
- [Search-R1](#): RL with reasoning and searching (tool-call) interleaved LLMs
- [Code-R1](#): Reproducing R1 for Code with Reliable Rewards
- [ReSearch](#): Learning to Reason with Search for LLMs via Reinforcement Learning
- [DeepRetrieval](#): Hacking Real Search Engines and retrievers with LLMs via RL for information
- [MetaSpatial](#): Reinforcing 3D Spatial Reasoning in VLMs for the Metaverse

Let's build together!